

Srinivas Sivakumar

Phone: +91 9869056567 Email: srinivas123.sivakumar@gmail.com
LinkedIn: linkedin.com/in/srinivas-siva GitHub: github.com/srinivassivakumar

Summary

AI/ML Engineer and Data Engineer with experience building production-focused AI agents, multimodal applications, predictive systems, and scalable data pipelines. Skilled in LLM integration, RAG, computer vision, FastAPI, PostgreSQL, AWS, Docker, Kubernetes, and MLflow. Experienced in developing validated AI workflows, deployment automation, and data-driven systems across manufacturing, networking, fashion, and wildlife conservation.

Skills

Languages: Python, TypeScript, SQL, Bash
AI and Machine Learning: PyTorch, TensorFlow, Hugging Face, LLMs, RAG, FAISS, CLIP, Diffusion Models, MLflow
Backend and Data: FastAPI, PostgreSQL, Apache Spark, Snowflake, Airflow, Dagster
Cloud and DevOps: AWS, Docker, Kubernetes, Jenkins, GitLab CI/CD, Railway, OpenStack, Linux, Infrastructure as Code

Experience

AI Software Engineer – Mexo Jun 2026 – Present
• As an early engineering hire on a two-person team at a very early-stage startup, independently built a manufacturing AI agent using TypeScript, Flue, OpenRouter, and PostgreSQL to analyze production progress, machine capacity, downtime, and operational dependencies, enabling early identification of delivery risks and evidence-based explanations.
• Built a Terraform-managed Amazon S3 backup infrastructure and migrated the company's database backups and historical data into centralized object storage.

R&D Engineer (AI) – Zybasis Dec 2025 – May 2026
• Developed an AI-based bandwidth prediction model to forecast network degradation and support proactive service-impact assessment.
• Built an automation system for customer-issue detection, root-cause analysis, and resolution workflows.
• Engineered a data pipeline integrating CheckMK and multiple infrastructure sources to process network and IOPS data for anomaly detection.

AI/ML Engineer (Freelance Engagement) – MyMirro Sep 2025 – Nov 2025
• Built an end-to-end multimodal AI fashion assistant that transforms wardrobe images into a personalized digital wardrobe, generates outfit combinations, provides styling recommendations, and suggests clothing purchases tailored to user preferences and occasions.
• Developed computer vision and RAG pipelines using Gemini API, FashionCLIP, CLIP, FAISS, SegFormer, and SAM to understand garments, retrieve relevant fashion knowledge, and generate context-aware outfit and shopping recommendations.
• Engineered and deployed the platform using FastAPI, MLflow, Docker, Terraform, and Railway, while building Python and TypeScript data pipelines with Playwright and Puppeteer to aggregate 1,000+ fashion products per run from 10+ e-commerce websites.

Data Engineer – Wildlife Conservation Trust Sep 2024 – Jul 2025
• Deployed and maintained RStudio Shiny servers with CI/CD automation using Jenkins, Airflow, and Dagster, enabling scientists to continuously access and visualize research data.
• Built AWS-based pipelines using S3, Lambda, Airflow, and Snowflake for real-time sensor data processing and delivery of Chambal River monitoring data.
• Developed image recognition systems with automated training and reporting pipelines to identify individual crocodiles during wildlife surveys.
• Created a GPS-free 3D tracking system using Python, AWS EC2, and a hybrid OpenStack deployment as a prototype for tracking underwater Gangetic dolphins.

Data Engineering Intern – NCBS (TIFR) Jun 2023 – Aug 2024
• Deployed Ceph using Docker with automated backup pipelines integrated into OpenStack.
• Built PySpark ETL pipelines using Databricks, Airflow, Snowflake, and Dagster.
• Deployed GitLab on OpenStack with automated backup and access-control systems.

Education and Certifications

• MCA – Storage and Cloud Computing, Jain University 2022 – 2024
Key Projects: DevOps pipeline for ML deployment and Infrastructure as Code automation.
• BCA – Computer Applications, Indira College, Pune University 2019 – 2022
Project: IoT-based methane detection sensor for food spoilage prevention.
• Cloud Big Data Engineer Certification, 8-Month Ultimate Big Data Masters Program – TrendyTech Insights Jul 2026
• Linux Administration Certification – Besant Technologies
• Patent: Design Patent for a light-based automatic feeder for aquarium fish.